

Which Pair Is It About?

Pair-Conditioned Verification of Literature-Mined Biotic Interactions

Anonymous ACL submission

Abstract

Literature-mining pipelines that populate biotic-interaction databases propose candidate (taxon, relation, taxon) triples by co-occurrence, then decide which to keep with a *sentence-level* classifier: does this passage describe an interaction? A candidate is retrieved precisely because its passage names two taxa near an interaction term, so that question is answered before it is asked; what needs deciding is whether the passage asserts an interaction between *these two*. We test the claim with a controlled comparison: one encoder, one 48k-row teacher-labelled corpus, three seeds, differing only in what the model is shown. On 437 expert-graded candidates, telling the model which pair the candidate is about raises AUPRC from 0.851 to 0.918 and F1 from 0.775 to 0.871 (75 items fixed against 22 broken, McNemar $p = 1.3 \times 10^{-7}$), at every one of 99 thresholds; adding the relation term to the query buys nothing further. The 32B model that labelled the corpus, asked the same questions zero-shot, shows the same gap at twice the size ($0.789 \rightarrow 0.943$). The same comparison on the public BioRED corpus, with candidates built the same way, gives $0.738 \rightarrow 0.843$ on the BioCreative VIII test set ($p = 5.9 \times 10^{-148}$) and +9.6 points of its standard entity-pair F1, and the gain lies where the argument puts it: none on sentences that name two entities, +0.120 on sentences that name more. We analyse the remaining errors, test candidate rules, and report the deployed model: a 110M-parameter pair-conditioned verifier that also labels direction in four categories, at 32 candidates per second on a CPU. Code: <https://anonymous.4open.science/r/XXXX>.

1 Introduction

Biotic-interaction databases such as GloBI (Poe-
len et al., 2014) are populated in part by text min-
ing. A retrieval system proposes candidate triples
(s_1, r, s_2) together with the passage that allegedly

supports them, and something must decide which
candidates are worth a curator’s attention. The natu-
ral choice, and the common one, is a sentence-level
interaction classifier: show it the passage and ask
whether it describes a biotic interaction.

That question is answered before it is asked. Can-
didates are proposed by co-occurrence: a passage
naming two taxa near an interaction term yields a
candidate whether or not it asserts that interaction
between those two. Every candidate therefore ar-
rives already satisfying, lexically, the property a
sentence-level classifier was trained to detect — in
the pipeline we study, on all 175,588 candidates
of its retrieval pool (§2). What actually needs de-
ciding is narrower: whether the passage asserts
an interaction between *these two* taxa. When a
passage names three organisms it can support one
candidate and refute the other two, and a classifier
that sees only the passage must give all three the
same answer.

We make this argument testable rather than
rhetorical. We train the same 110M-parameter en-
coder on the same teacher-labelled corpus three
times — reading the passage alone, the passage
with the candidate pair, and the passage with the
whole candidate triple — and compare them on a
437-item benchmark graded by a domain expert
(§4). Everything except the input is held fixed, so
differences between them are attributable to the
formulation and to nothing else. Telling the model
which pair it is judging wins by a wide and sig-
nificant margin; telling it the relation term as well
does not add to it. We then rebuild the setting on
BioRED (Luo et al., 2022), a public corpus with
gold entity and relation annotations, and the result
replicates there with the gain in the predicted place:
absent on sentences that name only two entities,
large on sentences that name more.

Contributions.

- A controlled comparison of sentence-level,

pair-conditioned and triple-conditioned verification on one encoder and one corpus, replicated on a public benchmark and in a zero-shot 32B model, with the gain located where the argument says it should be (§4).

- An account of what the pair query buys — disambiguation between the pairs a passage names, and a check on candidates whose arguments are not organisms — and a bound showing that checking a triple’s components separately cannot reach the largest error class (§4, §5).
- Deterministic candidate rules, validated on 42k training rows before evaluation, measured as a small precision gain rather than a large one (§6).
- A deployed configuration with a fixed threshold, four-way direction output and CPU throughput, and six negative results (§7, Appendix K).

2 Setting

Where candidates come from. BiotXplorer (Ruch et al., 2024) proposes candidate biotic-interaction triples from the literature indexed by SIB Literature Services (Gobeill et al., 2020). For each passage it matches taxon surface forms against a taxonomy and interaction surface forms against a vocabulary derived from the Relation Ontology, and emits every resulting (s_1, r, s_2) combination with the passage, the matched surface strings, and their canonical forms. This is the application we build for, and we use it only as the source of candidates: no component of it is a baseline in this paper.

The problem this paper addresses was diagnosed by BiotXplorer’s own authors. Ruch et al. (2024) graded 100 random BiotXplorer triples, found that the system identified the interacting species with a precision of 31%, and named the main cause of error: “instances where passages listed multiple species, which can be automatically filtered out.” We take the diagnosis as given and question only the remedy. On our benchmark, passages naming three or more taxa hold 154 of the 246 interacting pairs (63%), so filtering them out would discard most of what the pipeline exists to find. The alternative we test is to keep those passages and decide each candidate in them separately.

Why the sentence-level question is constant here. We pulled the retrieval pool in full: 175,588 can-

didate rows over 165,325 distinct passages and 115,404 distinct triplet keys. Both taxon surface forms and the interaction surface form are non-empty on all 175,588 rows, because the presence of an interaction term is the pool’s construction filter, and an independent lexical scan for GloBI interaction vocabulary fires on 99.83% of the same passages. A classifier trained to separate passages with interaction language from passages without it is, on this distribution, being asked about a constant. The missing input is not an annotator artefact but the query the pipeline itself issued.

Benchmark. One benchmark of 437 rows, 246 positive, in three blocks of different provenance; we report every block rather than pooling them into a number that hides their disagreement. BIOTX100 is retrieved candidates, each the top-ranked candidate for its passage, graded by a domain expert on four axes: whether each taxon surface form was correctly resolved, whether the relation was correct, and whether the whole triple is supported by its passage. Half were drawn from pairs already recorded in GloBI and half at random, and after removing three training duplicates it contributes 97 rows: 50 GloBI-matched with 43 positive, 47 random with 32 positive. It is therefore enriched for known interactors and for top-ranked candidates, which is why its positive rate is far above the 31% Ruch et al. (2024) measured on random triples at any rank. REJECT50 is candidates the retrieval pipeline discarded, curated for the ones it should have kept: 41 clean rows, 9 positive. It is the paper’s recall-side instrument (Appendix D). TEST299 is 299 pair-level items from expert curation, 0.542 positive, and is where passages most often name several taxa.

The blocks do not carry the same gold. BIOTX100 is the only block graded at the level of the whole triple; TEST299 was curated pair-level and its stored relation term was attached by retrieval rather than by an annotator. We therefore evaluate everywhere at the level all three share — *do these two taxa interact in this passage* — and use the triple-level grades only to diagnose (§5). This understates what triple conditioning buys, since a candidate with the right pair and the wrong relation counts as correct, and we prefer understating it to scoring two blocks against a relation term nobody checked.

Query representation. Each candidate carries a canonical string and the surface form matched

in the passage. On BIOTX100 these differ for s_1 , s_2 and r on 34, 38 and 59 of 100 rows, and the canonical string is absent from the passage on 31, 27 and 54. We always query with the surface form, which is what the model can see.

Evaluation protocol. Thresholds are where this kind of evaluation goes wrong most easily, so the policy is fixed in advance. A model’s reported operating point is either **pre-specified** — chosen on its own training-development split, never on the benchmark — or **block-held-out**: fitted on two of the three source blocks and applied to the third, so that every row is scored by a threshold that never saw it and every row is used. Comparisons between models use the same policy for both. Intervals are Clopper–Pearson; McNemar tests (McNemar, 1947) carry a continuity correction. Where three seeds exist we report the per-seed mean and standard deviation, and where we report the average of three checkpoints’ probabilities we call it an ensemble. The benchmark ships an `in_train` flag on 9 rows; a 5-gram Jaccard scan against all 48,338 training passages finds 3 more at Jaccard 1.000, all in BIOTX100. All 12 are excluded throughout.

The gold is pair-level and therefore order-free: *do these two taxa interact*, in either direction. A model that reads the candidate in order is scored the same way, as the higher of its two scores over the two argument orders — the probability that the pair interacts in at least one direction. We fix this before scoring because the alternative penalises a model for rejecting a correctly paired candidate whose stated relation runs backwards, which the gold does not count as an error. The sentence-level model reads no candidate and is order-free already.

Two things about this benchmark’s history should be on the record. It is not untouched: during the development of the verifier it was scored for roughly a hundred model configurations, so the comparisons in this paper are between recipes chosen with it in view, and the controlled comparison of §4 is the result we would expect to survive that. And one gold label has been changed since the benchmark was frozen: a TEST299 candidate pairing *Acanthamoeba* with *Pseudomonas*, revised from positive to negative after the curator, annotating direction without model scores, recorded that the two co-occur in a shared host rather than interacting.

3 Two formulations, one corpus

The teacher. Labels come from a locally hosted 32B open-weight instruction model (Qwen Team, 2025), prompted zero-shot at temperature 0 and asked, per candidate, whether the passage supports *this* triple. Its labels therefore vary across candidates that share a passage. Against 150 human-graded items its mislabel rate is $22/150 = 14.7\%$, Clopper–Pearson [9.4%, 21.4%], split 10 false positives and 12 false negatives.

The corpus. 48,338 teacher-labelled rows, 0.343 positive, drawn from the same retrieval pool as the benchmark (disjoint passages): a 20,000-row base sample, 16,546 *pair-contrast* rows, 6,201 actively sampled rows, and 5,591 rows with the candidate’s argument order reversed. The pair-contrast block is what gives the corpus its point: 2,017 passages enter as several rows that are identical in the passage and differ only in the candidate and the label.

Three models, one recipe. All three are cross-encoders on BiomedBERT-base (Gu et al., 2021) (microsoft/BiomedNLP-BiomedBERT-base-uncased-110M parameters), distilled from the teacher’s labels (Buciluă et al., 2006; Hinton et al., 2015) with the same recipe: AdamW at learning rate 2×10^{-5} and weight decay 0.01, a linear schedule with 10% warm-up, batch size 32, inputs of at most 256 wordpieces, three epochs, gradient clipping at 1.0, and the checkpoint with the best development AUPRC, on a 10% development split grouped by taxon pair so that no pair spans training and development. Three seeds each; one run takes about four minutes on a single A100. They differ only in their input. The **sentence-level baseline** reads the passage alone. It is the formulation a pipeline reaches for first, and on rows that share a passage it is shown the same input with different labels, which it cannot fit. The **pair-conditioned** model reads the two candidate taxa, s_1 [SEP] s_2 , as segment A and the passage as segment B (Nogueira and Cho, 2019). The **triple-conditioned verifier** reads the whole candidate, s_1 [SEP] r [SEP] s_2 . The middle model separates what the pair buys from what the relation adds.

4 Results

4.1 Conditioning on the pair

Table 1 gives the comparison. Telling the model which pair a candidate is about raises per-seed AUPRC from 0.851 ± 0.006 to 0.918 ± 0.005

System	AUPRC	P	R	F1
Accept everything	0.563	0.563	1.000	0.720
Sentence-level	0.851 $\pm$.006	0.730	0.825	0.775
Pair-conditioned	0.918$\pm$.005	0.849	0.894	0.871
Triple-conditioned	0.909 $\pm$.004	0.820	0.886	0.852
<i>The 32B teacher, zero-shot, asked the same three questions</i>				
Sentence question	0.789	0.640	0.947	0.764
Pair question	0.943	0.894	0.858	0.876
Triple question	0.946	0.855	0.886	0.870

Table 1: The controlled comparison on 437 rows: one encoder, one corpus, only the input differs. AUPRC is the per-seed mean $\pm$ standard deviation over three seeds; P/R/F1 are three-checkpoint ensembles at block-held-out thresholds. The teacher’s score is $P(\text{YES})$ from its first answer token; its P/R/F1 use the same block-held-out thresholds, and the triple question is its labelling prompt, asked in both argument orders and maxed like the trained triple arm (prompts in Appendix E). Against the sentence-level model, the pair-conditioned one fixes 75 items and breaks 22 (McNemar $p = 1.3 \times 10^{-7}$) and the triple-conditioned one fixes 74 and breaks 32 ($p = 6.8 \times 10^{-5}$); triple against pair, 20 fixed and 31 broken ($p = 0.16$). Per-seed values are in Appendix A.

and, at block-held-out thresholds, F1 from 0.775 to 0.871, precision and recall both. On decision correctness the pair-conditioned model fixes 75 items and breaks 22, McNemar $p = 1.3 \times 10^{-7}$. The advantage does not depend on where we cut: it is ahead at every one of 99 thresholds, by at least 0.058 F1 (Figure 1).

Adding the relation term to the query does not add to this. The triple-conditioned model is also far ahead of the sentence-level one ($p = 6.8 \times 10^{-5}$), but against the pair-conditioned model it fixes 20 items and breaks 31 ($p = 0.16$), and its AUPRC is, if anything, lower. What the decision needs is the pair. The one place the relation term earns its keep is at strict operating points: at $\tau = 0.5$ both reach precision 0.913, the triple model at recall 0.768 and the pair model at 0.683 (Table 5).

The same gap in the teacher. The formulation effect is not a property of small encoders. Asked the sentence-level question zero-shot, the 32B teacher scores AUPRC 0.789, below the 110M sentence-level student; asked about the pair, 0.943 (Table 1). Its gap, +0.155, is twice the student’s, and it sits in the same place: +0.180 on passages naming three or more taxa against +0.104 on the rest, where the student’s is +0.084 against +0.042. A model three hundred times smaller that is asked the right question beats the teacher asked the wrong one by 0.13 AUPRC. As with the stu-

Candidates	n	Sentence	Pair	Triple
passage names ≤ 2 taxa	161	0.876	0.918	0.879
passage names ≥ 3 taxa	276	0.846	0.930	0.930
both arguments are taxa, ≤ 2	33	0.945	0.970	—
both arguments are taxa, ≥ 3	227	0.878	0.947	—
an argument is not a taxon, ≤ 2	128	0.856	0.908	—
an argument is not a taxon, ≥ 3	49	0.628	0.808	—
TEST299	299	0.819	0.922	0.915
BIOTX100	97	0.954	0.970	0.938
REJECT50	41	0.465	0.414	0.503

Table 2: Where the gain lives, AUPRC of the three ensembles. Taxa are counted as distinct mentions found in the passage by TaxoNERD (Le Guillaume and Thuiller, 2022) (en_ner_eco_biobert); an argument “is not a taxon” when TaxoNERD does not recognise it in the passage. TEST299 passages name 5.1 taxa on average, BIOTX100 2.6 and REJECT50 1.7.

dent, the relation term adds nothing (0.946 against 0.943). Nor is the gap a property of one model size: from Qwen3-1.7B to Qwen3.5-122B the pair question leads, by +0.04 to +0.20 AUPRC on our benchmark and +0.11 to +0.32 on BioRED (Appendix F). Against its own teacher the pair-conditioned student gives up 0.027 AUPRC, which distilling the teacher’s probabilities does not recover (Appendix B), and nothing in F1 at block-held-out thresholds (0.871 against 0.870), at 32 candidates per second on a CPU.

Table 2 locates the gain, and it is the most direct test of the argument. On passages that name three or more taxa, the pair query is worth +0.084 AUPRC. On passages naming two or fewer it is worth less, +0.042, and the split below shows why it is worth anything there at all: the candidates come from automatic matching, and in 41% of them one “taxon” is a string TaxoNERD does not recognise as an organism in the passage — an organ, an author’s name, a vernacular term. Where both arguments are recognised taxa and the passage names no third one, the case in which the sentence-level question and the pair question coincide, the pair query adds +0.026 on 33 items; where an argument is not a taxon it adds +0.052, and +0.180 when the passage also names others. The pair query does two things, then: it tells the model which of several pairs to judge, and it lets the model check that what it is judging is a pair of organisms. BioRED, whose entities are gold-annotated, isolates the first.

Model	AUPRC, sentences naming two concepts				Query given to the triple model	P	R	F1	AUPRC
	AUPRC	F1	2 concepts	≥ 3					
Sentence-level	0.738 $\pm$.004	0.718	0.909	0.727	Retrieved triple	0.820	0.886	0.852	0.909
Pair-conditioned	0.843 $\pm$.004	0.817	0.902	0.847	Generic relation, pair kept	0.897	0.671	0.767	0.888
+ entity markers	0.840 $\pm$.003	0.815	0.907	0.844	Relation shuffled across items	0.881	0.720	0.792	0.904
					Relation only, no taxa	0.833	0.020	0.040	0.677
					Empty query	1.000	0.037	0.071	0.641
<i>Zero-shot 32B teacher, 3,000-candidate sample</i>									
Sentence question	0.566	0.662	0.921	0.555	Table 4: Degrading segment A of the trained triple-conditioned model at inference, weights and its own block-held-out thresholds unchanged. Base rate 0.563. Per-seed standard deviations on AUPRC are ≤ 0.027 .				
Pair question	0.813	0.766	0.958	0.799					

Table 3: BioRED, BioCreative VIII test set: 16,606 sentence-level candidates from 400 abstracts, 0.463 positive, 538 from sentences naming two concepts. AUPRC is the per-seed mean $\pm$ standard deviation over three seeds; F1 is at the pre-specified threshold (mean of the checkpoints’ development thresholds), the teacher’s at its greedy answer. EP F1 is BioRED’s document-level entity-pair F1, a pair scoring the maximum over the sentences that co-mention it. On the teacher’s sample the trained models score 0.757 and 0.860. Pair against sentence-level: 2,642 items fixed, 1,064 broken (McNemar $p = 5.9 \times 10^{-148}$).

4.2 A public benchmark: BioRED

The biodiversity benchmark is ours, singly annotated, and modest in size. To test the claim where anyone can check it, we rebuild the same setting on BioRED (Luo et al., 2022), a public corpus of PubMed abstracts with expert-annotated entities and relations. Candidates are generated exactly as the biodiversity pipeline generates them, by co-occurrence: every pair of distinct annotated concepts co-mentioned in one sentence, restricted to the eight entity-type pairs that carry relations in training, with the sentence as passage. A candidate is positive when BioRED relates the two concepts anywhere in the abstract, a document-level label whose noise every model shares. We follow the protocol of Lai et al. (2025): models train on the 500 abstracts of BioRED’s training and development splits (15,640 candidates), select their epoch and threshold on its 100 test abstracts, and are scored on the 400 abstracts of the BioCreative VIII test set (Islamaj et al., 2024). BioRED has no retrieved relation term, so only the sentence-level and pair-conditioned formulations are compared, trained with the recipe of §3.

The result (Table 3) reproduces the biodiversity one in both its parts. Conditioning on the pair raises AUPRC from 0.738 to 0.843 and F1 from 0.718 to 0.817. And the gain sits in the same place: on candidates from sentences that name only two concepts the two models are indistinguishable (0.909 against 0.902; per seed 0.899–0.919 against 0.892–0.914),

while on those from sentences naming more, which are 97% of the candidates, conditioning is worth +0.120. Marking the two mentions in the passage, the standard relation extraction input (Zhou and Chen, 2022), adds nothing to the pair query ($p = 0.08$). The teacher, asked the sentence-level question, accepts 92% of its sample (AUPRC 0.566 at a base rate of 0.48); asked about the pair, 0.813, still behind the trained pair-conditioned model on these document-level labels (0.860 on the same sample).

On BioRED’s standard document-level metric the pair-conditioned verifier reaches an entity-pair F1 of 67.1, against 57.5 for the sentence-level one. Systems that read the whole abstract reach 74.1 to 75.3 (PubMedBERT, BioREx and BioREDirect; Lai et al., 2023, 2025) and zero-shot GPT-4 42.9 (Islamaj et al., 2024). Much of that gap is structural: a quarter of the related pairs are never co-mentioned in one sentence, which caps any sentence-scope verifier at recall 0.749. Given the whole abstract, the same model and recipe reach 67.0 (recall 0.697, precision 0.646); what the published systems add is document-level modelling, not the question we ask.

4.3 The query, not the encoder

Degrading the trained triple model’s query at inference, with weights and thresholds unchanged, tells the same story from the other side (Table 4). Deleting both taxa collapses ranking from 0.909 to 0.677 against a base rate of 0.563, and emptying the query entirely gives 0.641. Replacing or shuffling the relation, pair kept, barely moves ranking (0.888, 0.904): the model leans on the pair, which is what the trained pair-only model already showed.

The empty-query arm is *not* the sentence-level baseline: deprived of the query it was trained with, it does far worse (0.641) than a model trained without one (0.851); hence three trained arms.

Nor does the gain come from the corpus’s pair-




Figure 1: Precision, recall and F1 of the pair-conditioned model against the decision threshold on all 437 rows, with the sentence-level model’s F1 for comparison (same corpus and encoder). The shaded band spans the three block-held-out thresholds. The pair model is ahead at every τ .

contrast design: a triple-conditioned model trained on the 20,000-row base sample alone reaches AUPRC 0.918 ± 0.003 (full corpus 0.909 ± 0.004) and the pair model’s block-held-out F1, 0.871.

Nor from the encoder. Besides the recipe of §3 we tried seven others on development data only: five epochs at learning rates 1, 2 and 3×10^{-5} , BiomedBERT-large at 1 and 2×10^{-5} , and BioLinkBERT-base and -large (Yasunaga et al., 2022). The three ahead of it by 0.01 or more on both development splits were scored once on test (Appendix C). None pays on both, so we keep the base recipe: BiomedBERT-large moves neither BioRED model, BioLinkBERT-base lifts the BioRED pair model (AUPRC 0.855, EP F1 67.9) but not ours (F1 0.861 against 0.871), and BioLinkBERT-large lifts both a little, neither significantly. With every encoder trained in both formulations the pair query wins: +0.103 to +0.113 AUPRC on BioRED, +0.052 to +0.069 on ours.

4.4 An operating curve, not a verdict

The pair model’s precision moves from 0.745 to 0.963 as τ moves (Table 5), and the right point on that curve is not a property of the model. Ingestion into a reference database pays for false positives and belongs at the high end; assembling candidates for expert review pays for false negatives and belongs at the low end. A filter that emits a verdict fixes that choice at training time, for users who were never asked. A calibrated score with an ad-

τ	sentence			pair			triple		
	P	R	F1	P	R	F1	P	R	F1
0.01	.690	.915	.787	.745	.976	.845	.721	.935	.814
0.04	.763	.813	.787	.795	.943	.862	.790	.919	.850
0.10	.796	.732	.763	.835	.907	.869	.823	.907	.863
0.25	.837	.687	.754	.889	.850	.869	.877	.841	.859
0.50	.866	.524	.653	.913	.683	.781	.913	.768	.834
0.90	.930	.163	.277	.963	.423	.588	.945	.626	.753

Table 5: The three models’ operating curves on all 437 rows, three-checkpoint ensembles. Rows show the shape; the reported operating points are those of Table 1.

justable threshold leaves it to them (Geifman and El-Yaniv, 2017).

5 What the remaining errors are made of

The pair-conditioned model’s 39 false positives at its block-held-out operating point are not homogeneous: 27 come from passages naming three or more taxa, and 17 have an argument TaxoNERD does not recognise as a taxon. Read by hand, they sort into a small number of mechanisms (examples in Table 12, Appendix H; the per-row table is released): *pair-binding*, where both taxa and the relation are correct but the relation’s argument is a third organism; *co-listing*, where the two taxa are members of one enumeration and interact with neither each other nor anything in particular; *negation*, where the passage asserts the opposite; *same organism*, where the two “taxa” are one organism named twice or an organism and its own clade; candidates whose relation term names no biotic interaction at all (*relative to*, *hybridization*, *migration*); and *entity* failures, where a taxon string is a taxonomic author, a place, an organ, or an adjective inside a pathogen’s name (“*equine* influenza virus”, which the pair-conditioned model already rejects). A further handful read as gold errors and are discussed in the Limitations.

A bound on component-wise verification. The pair-binding class has a structural property worth stating, because it rules out the obvious alternative architecture. Its members have both taxa correctly resolved *and* the correct relation; what is wrong is a binding that exists only in retrieval’s combinatorics. A verifier that checks s_1 , r and s_2 separately — however good each check is — cannot see it. On the 97 clean BIOTX100 candidates, where triple-level grades make the classes countable, 15 of the 32 unsupported candidates are pair-binding failures (the rest: 9 entity, 8 predicate), a share of 0.47,

Clopper–Pearson $[0.29, 0.65]$. A component-wise verifier accepting a fraction r of the 65 supported triples therefore also accepts all 15, so its precision is at most $65r/(65r + 15)$: 0.81 at full recall, $[0.76, 0.88]$ across that interval. Only a decision that reads the pair *in context* can go above it.

6 Candidate rules

Several of these mechanisms are questions about the candidate rather than the passage — whether two strings name one organism, whether a word is a surname, whether “relative to” is an interaction — and a classifier is a poor instrument for them. We answer them with eight deterministic rejection rules applied before the verifier decides (Appendix I). No rule has a parameter fitted to evaluation labels, and since four were written after reading the false positives above, each had to pass a criterion fixed in advance on the 42,747 training rows (at least 85% of what it rejects is teacher-negative); they were then frozen, recorded by hash and evaluated once. The result (Table 13) is smaller than the motivation suggested, and the reason is the headline of this paper. On the pair-conditioned model the rules buy 0.009 to 0.018 precision for 0.008 to 0.020 recall, F1 unchanged ($p = 1.00$): of its 39 false positives they catch 6, because a model told which pair it is judging already rejects most of what they were written to reject. On the triple-conditioned model, which leans on its relation term, they do more (precision $0.820 \rightarrow 0.862$, 14 fixed and 5 broken, $p = 0.067$). Both are upper estimates, since half the rules were chosen with this benchmark’s errors in view; their case in deployment is transparency, since each rejection names the rule that made it. A taxonomy gate is no substitute: surnames, clades and transposon families all resolve in a taxonomy (Appendix J).

7 The deployed configuration

In the pipeline, verification runs as one 110M-parameter model with two heads. It differs from the pair-conditioned model of §3 in three respects, each of which we adopted during development: the two taxa are given in alphabetical order, so the input is identical whichever way round retrieval presents the pair; they are marked in place in the passage ($@ \dots @$, $\# \dots \#$) rather than only named in segment A; and the training labels are the teacher’s pair-level judgments. Its interaction threshold, 0.5, was set before evaluation and never tuned. On the

Expert direction label	n	answered	correct
forward or reverse	84	62 (74%)	55 of 62 (0.887)
undecidable (“?”)	12	—	9 left UNCERTAIN
bidirectional	1	—	1 BIDIRECTIONAL

Table 6: The deployed model’s four-way direction output against 97 expert direction labels. “Answered” counts FORWARD or REVERSE; the rest are UNCERTAIN. The abstention level (0.60) was chosen by the expert after seeing this set’s coverage curve, so the figures are optimistic.

437 rows it reaches AUPRC 0.934, and precision 0.851, recall 0.955, F1 0.900; with the candidate rules, precision 0.869, recall 0.947, F1 0.907 (6 items fixed, 2 broken, $p = 0.29$).¹ It needs no GPU: on eight CPU threads it scores 32.3 candidates per second with both heads, faster than the pipeline retrieves them. Where precision matters more than speed, its least confident candidates can go to the zero-shot teacher: escalating the 30% nearest the threshold to Qwen3.5-122B gives precision 0.941 at recall 0.911, a precision the model alone reaches only at $\tau = 0.99$, with recall 0.740 (Appendix G).

The second head reports the interaction’s *direction* in four categories. FORWARD and REVERSE say which taxon acts on the other; they come from a direction score that is exactly antisymmetric in the two marked taxa, and are reported only when it is confident, otherwise the answer is UNCERTAIN. BIDIRECTIONAL is assigned when the relation is mutual by definition — mutualism, symbiosis, competition, co-infection, *interacts with* — as listed in a relation-polarity lexicon derived from the Relation Ontology; for those, neither taxon is the subject and the model is not asked. Expert curation of the benchmark scores reversed triples as correct (Appendix J), so direction is evaluated on a separate expert annotation of 97 candidates (Table 6): the model answers 74% of the directional ones, correctly on 0.887 of those.

8 Related work

Classifiers that answer an easier question. That a classifier can score well by exploiting how its data were assembled is among the best-established results in NLP. Partial-input baselines recover much of the label from the hypothesis alone in natu-

¹The *co_listed* rule is omitted from this configuration: on the deployed model’s accept set it removes one false positive at the cost of three true ones.

ral language inference (Gururangan et al., 2018; Poliak et al., 2018) and from the passage alone in reading comprehension (Kaushik and Lipton, 2018), and such rules hold in distribution while failing under deployment (Geirhos et al., 2020; Du et al., 2024); the closest task-level precedent is fact verification (Thorne et al., 2018), where a claim-only classifier matches evidence-aware models because of how the corpus was built (Schuster et al., 2019). Our sentence-level baseline is a partial-input model in exactly this sense, and the specific feature here is that the missing input is the query the pipeline itself issued rather than an annotator artefact (Wiegand et al., 2019; Parmar et al., 2023). Feng et al. (2019) caution that partial-input results license claims about models and corpora rather than about tasks, and we keep ours to the former.

The argument-ignoring failure is already named.

Distant supervision assumes that a sentence containing two related entities expresses the relation (Mintz et al., 2009), and its standard relaxation exists because it does not (Riedel et al., 2010). Since TACRED an instance has been the triple (sentence, e_1 , e_2) rather than the sentence (Zhang et al., 2017), and how the pair is injected into the encoder is worth several F1 points (Baldini Soares et al., 2019; Wu and He, 2019; Zhou and Chen, 2022). Rosenman et al. (2020) name our failure exactly, as an “event heuristic” that classifies on (sentence, relation) while ignoring the arguments. *We claim no part of that diagnosis.* What is ours is a controlled measurement of what it costs in a verification setting built by co-occurrence retrieval, located by block; a bound on component-wise remedies; and the candidate rules. The entity-type shortcut usually invoked for such errors (Peng et al., 2020; Mtumbuka and Schockaert, 2024) is unavailable here, since both slots are filled from one inventory of taxon surface forms. Error attribution across pipeline stages is standard and easy to get wrong (Taillé et al., 2020; Ai and Kavuluru, 2023).

Conditioning on the query. Casting a relation as a question about a named entity (Levy et al., 2017), or as an entailment hypothesis verbalising the relation between two of them (Obamuyide and Vlachos, 2018; Sainz et al., 2021), is the reformulation our input format instantiates, and the resulting decision — does this passage support this specific claim — has the shape of scientific claim verification (Wadden et al., 2020), and of validating a knowledge-graph triple against text (Lehmann

et al., 2012), with the difference that those systems retrieve evidence for a given triple while here the passage is fixed by the retrieval that proposed the triple and the question is whether that passage, and no other, supports it. Distilling a large model’s judgments into a small student is established practice (West et al., 2022; Gekhman et al., 2023). Selective classification thresholds a score instead of issuing a verdict (Geifman and El-Yaniv, 2017), which is how extraction is made useful to human abstractors: Swaminathan et al. (2024) abstain under calibrated thresholds to automate 57–95% of clinical chart abstraction at controlled error.

Text mining for interaction databases. GloBI (Poelen et al., 2014) is populated in part by text mining (Dimitrova et al., 2020). Judging an extracted fact before it reaches a curator is an old idea (Rodriguez-Esteban et al., 2006), as is evaluating such tools by whether they help the curator (Hirschman et al., 2012; Britan et al., 2018; Lee et al., 2018). Cuzick et al. (2023), independently, take the interacting pair rather than the passage as the annotation unit — the unit our pair-binding class isolates. Large language models are applied to interaction extraction at scale (Keck et al., 2025; Farrell et al., 2024; D’Souza et al., 2025; Zou et al., 2026), but as extractors: Zou et al. (2026) standardise names and categories post hoc without checking extracted interactions against the text that produced them. We are not aware of published work that verifies literature-mined biotic-interaction triples against their source passage before ingestion.

9 Conclusion

When candidates are retrieved by co-occurrence, the sentence-level question is answered before it is asked, and a classifier built for it is left with little to decide. Telling the model which pair the candidate is about — with the same encoder and the same data — raises AUPRC from 0.851 to 0.918 on the biodiversity benchmark and from 0.738 to 0.843 on BioRED, the same question lifts a zero-shot LLM at every size from 1.7B to 122B parameters, and the relation term adds nothing on top. The gain sits where the argument puts it: on passages that name more than one candidate pair, and, where candidates come from automatic matching, on candidates whose arguments are not organisms. The result is a calibrated score where a verdict stood, with a four-way direction label beside it, cheap enough to run on every candidate.

Limitations

One benchmark, singly annotated. 437 rows in three blocks, every label assigned by one domain expert, with no independent second grading; we report no inter-annotator agreement and cannot bound label noise. Some apparent false positives are not model errors in the usual sense. A passage stating that “the bat tick . . . has been reported . . . to occasionally feed on humans” is labelled negative, correctly, because the curator judged the candidate’s *resolved* taxon, which retrieval had mapped to the wrong species; the model sees only the surface form, which the passage does support. Errors of this kind are upstream, and no passage-reading model can repair them. Candidates on which two independent models confidently contradict the gold go to blind expert re-annotation, mixed with an equal number of random controls and with labels and scores hidden; no label is changed on the strength of a model’s disagreement. On the recall side the denominator is 9 positives.

Top-ranked candidates only. Every BIOTX100 row is the top-ranked candidate for its passage, and half are pairs GloBI already records. Whether the gains hold further down the retrieval ranking, where candidates are noisier, is not measured here; TEST299 and the public benchmark of §4.2 are not rank-restricted.

The rules are partly fitted by choice. Four of the eight rules were selected after reading this benchmark’s false positives (§6). The training-data criterion bounds how much that selection can have memorised the benchmark, but it does not remove it, and their effects should be read as optimistic until confirmed on candidates neither the rules nor their author have seen. They fire on no row of a separate 100-candidate sample, so that sample can neither confirm nor contradict them.

Fresh candidates are harder. On that sample — 100 retrieved candidates drawn after the benchmark, half of them from passages containing several candidate pairs — all three models score well below their benchmark figures, though in the same order: AUPRC 0.665 sentence-level, 0.785 pair-conditioned, 0.765 triple-conditioned. Its label provenance is not identical to the benchmark’s, so we do not read the drop as measured, but it is the clearest sign that multi-pair passages remain the hard case and that the benchmark may understate their share of a live pool.

The diagnosis is triple-level; the evaluation is pair-level. Only BIOTX100 carries whole-triple gold. Evaluating at the pair level everywhere means a candidate with the right pair and a misbound relation counts as correct, so the comparison cannot credit the relation term with the cases it alone could catch; triple-level gold on the other two blocks is what would test whether the triple model’s advantage at strict operating points is real.

BioRED’s labels are document-level. A BioRED candidate is positive when the abstract relates the two concepts anywhere, so a co-mention in a sentence that does not state the relation counts as positive. Both models are trained and scored on the same noisy labels, so the comparison is controlled, but the absolute figures understate what either would reach against sentence-level gold.

Teacher noise and pretraining exposure. The training labels come from a prompted teacher wrong about one time in seven, and the student inherits that boundary; training on model-generated labels can degrade encoders (Lu and Smith, 2025). The passages are published MEDLINE text and the teacher is a pretrained model, so we cannot exclude that it saw them, and its labels may be optimistic for genuinely new literature.

Upstream failures are out of reach. Of the 50 candidates the retrieval pipeline discarded, 32 carry an incorrect interaction term and 17 a mis-resolved taxon, 14 of them both. Those fail before any verifier is consulted, and no change to the decision stage, ours included, addresses them.

One system, one domain. One retrieval system, one interaction vocabulary, taxon–taxon interactions in MEDLINE. The mechanism — a pair-agnostic decision over passages that name several candidate pairs — is not specific to ours, but we have not shown that other pools fail the same way.

References

- Xuguang Ai and Ramakanth Kavuluru. 2023. End-to-end models for chemical–protein interaction extraction: Better tokenization and span-based pipeline strategies. In *2023 IEEE 11th International Conference on Healthcare Informatics (ICHI)*, pages 610–618. IEEE. ArXiv:2304.01344.
- Livio Baldini Soares, Nicholas FitzGerald, Jeffrey Ling, and Tom Kwiatkowski. 2019. [Matching the blanks](#):

766	Distributional similarity for relation learning. In <i>Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics</i> , pages 2895–2905, Florence, Italy. Association for Computational Linguistics.	821
767		822
768		823
769		824
770		825
771	Aurore Britan, Isabelle Cusin, Valérie Hinard, Luc Mot- tin, Emilie Pasche, Julien Gobeill, Valentine Rech de Laval, Anne Gleizes, Daniel Teixeira, Pierre-André Michel, Patrick Ruch, and Pascale Gaudet. 2018. Ac- celerating annotation of articles via automated ap- proaches: evaluation of the neXtA5 curation-support tool by neXtProt . <i>Database</i> , 2018:bay129.	826
772		827
773		828
774		829
775		830
776		831
777		832
778	Cristian Bucilua, Rich Caruana, and Alexandru Niculescu-Mizil. 2006. Model compression. In <i>Pro- ceedings of the 12th ACM SIGKDD International Conference on Knowledge Discovery and Data Min- ing (KDD)</i> , pages 535–541.	833
779		834
780		835
781		836
782		
783	Alayne Cuzick, James Seager, Valerie Wood, Martin Ur- ban, Kim Rutherford, and Kim E. Hammond-Kosack. 2023. A framework for community curation of inter- species interactions literature . <i>eLife</i> , 12:e84658.	837
784		838
785		839
786		840
787		841
788		842
789		
790	Mariya Dimitrova, Jorrit Poelen, Georgi Zhelezov, Teodor Georgiev, Donat Agosti, and Lyubomir Penev. 2020. Semantic publishing enables text mining of biotic interactions . <i>Biodiversity Information Science and Standards</i> , 4:e59036.	843
791		844
792		845
793		846
794		847
795		848
796		
797	Jennifer D’Souza, Zachary Laubach, Tarek Al Mustafa, Sina Zarriß, Robert Frühstückl, and Phyllis Illari. 2025. Mining for species, locations, habitats, and ecosystems from scientific papers in invasion biology: A large-scale exploratory study with large language models . In <i>Proceedings of the 1st Workshop on Ecol- ogy, Environment, and Natural Language Processing (NLP4Ecology2025)</i> , pages 16–23, Tallinn, Estonia. University of Tartu Library.	849
798		850
799		851
800		852
801		853
802		854
803		855
804		856
805		857
806		
807	Maxwell J. Farrell, Nicolas Le Guillarme, Liam Brierley, Bronwen Hunter, Daan Scheepens, Anna Willoughby, Andrew Yates, and Nicole Mideo. 2024. The chang- ing landscape of text mining: a review of approaches for ecology and evolution . <i>Proceedings of the Royal Society B: Biological Sciences</i> , 291(2027):20240423.	858
808		859
809		860
810		861
811		
812	Shi Feng, Eric Wallace, and Jordan Boyd-Graber. 2019. Misleading failures of partial-input baselines . In <i>Proceedings of the 57th Annual Meeting of the Asso- ciation for Computational Linguistics</i> , pages 5533– 5538, Florence, Italy. Association for Computational Linguistics.	862
813		863
814		864
815		865
816		866
817		867
818		868
819		869
820		
	Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. 2020. Shortcut learning in deep neural networks . <i>Nature Machine Intelligence</i> , 2(11):665–673.	870
		871
		872
	Zorik Gekhman, Jonathan Herzig, Roei Aharoni, Chen Elkind, and Idan Szpektor. 2023. TrueTeacher: Learning factual consistency evaluation with large language models . In <i>Proceedings of the 2023 Con- ference on Empirical Methods in Natural Language Processing</i> , pages 2053–2070, Singapore. Associa- tion for Computational Linguistics.	873
		874
		875
		876
		877
	Martin Gerner, Goran Nenadic, and Casey M. Bergman. 2010. LINNAEUS: A species name identification system for biomedical literature . <i>BMC Bioinformat- ics</i> , 11:85.	
	Julien Gobeill, Déborah Caucheteur, Pierre-André Michel, Luc Mottin, Emilie Pasche, and Patrick Ruch. 2020. SIB literature services: RESTful customizable search engines in biomedical literature, enriched with automatically mapped biomedical concepts . <i>Nucleic Acids Research</i> , 48(W1):W12–W16.	
	Yu Gu, Robert Tinn, Hao Cheng, Michael Lucas, Naoto Usuyama, Xiaodong Liu, Tristan Naumann, Jianfeng Gao, and Hoifung Poon. 2021. Domain-specific lan- guage model pretraining for biomedical natural lan- guage processing. <i>ACM Transactions on Computing for Healthcare</i> , 3(1):1–23.	
	Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. 2018. Annotation artifacts in natural lan- guage inference data . In <i>Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)</i> , pages 107–112, New Orleans, Louisiana. Associa- tion for Computational Linguistics.	
	Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015. Distilling the knowledge in a neural network. In <i>NIPS Deep Learning and Representation Learning Workshop</i> .	
	Lynette Hirschman, Gully A. P. C. Burns, Martin Krallinger, Cecilia Arighi, K. Bretonnel Cohen, Alfonso Valencia, Cathy H. Wu, Andrew Chatr- Aryamontri, Karen G. Dowell, Eva Huala, Anália Lourenço, Robert Nash, Anne-Lise Veuthey, Thomas Wiegiers, and Andrew G. Winter. 2012. Text mining for the biocuration workflow . <i>Database</i> , 2012:bas020.	
	Rezarta Islamaj and 1 others. 2024. The overview of the BioRED (Biomedical Relation Extraction Dataset) track at BioCreative VIII. <i>Database</i> , 2024.	
	Divyansh Kaushik and Zachary C. Lipton. 2018. How much reading does reading comprehension require? A critical investigation of popular benchmarks . In <i>Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing</i> , pages	

878	5010–5015, Brussels, Belgium. Association for Com-	Frank Martin Mtumbuka and Steven Schockaert. 2024.	931
879	putational Linguistics.	Entity or relation embeddings? an analysis of encod-	932
		ing strategies for relation extraction . In <i>Findings</i>	933
880	François Keck, Henry Broadbent, and Florian Altermatt.	<i>of the Association for Computational Linguistics:</i>	934
881	2025. Extracting massive ecological data on state and	<i>EMNLP 2024</i> , pages 6003–6022, Miami, Florida,	935
882	interactions of species using large language models .	USA. Association for Computational Linguistics.	936
883	<i>bioRxiv</i> . Preprint, posted 27 January 2025.		
		Rodrigo Nogueira and Kyunghyun Cho. 2019. Pas-	937
884	Po-Ting Lai, Chih-Hsuan Wei, Ling Luo, Qingyu	sage re-ranking with BERT . <i>arXiv preprint</i>	938
885	Chen, and Zhiyong Lu. 2023. BioREx: Improving	<i>arXiv:1901.04085</i> .	939
886	biomedical relation extraction by leveraging hetero-		
887	geneous datasets. <i>Journal of Biomedical Informatics</i> ,	Abiola Obamuyide and Andreas Vlachos. 2018. Zero-	940
888	146:104487.	shot relation classification as textual entailment . In	941
		<i>Proceedings of the First Workshop on Fact Extraction</i>	942
889	Po-Ting Lai, Chih-Hsuan Wei, Shubo Tian, Robert Lea-	<i>and VERification (FEVER)</i> , pages 72–78, Brussels,	943
890	man, and Zhiyong Lu. 2025. Enhancing biomedical	Belgium. Association for Computational Linguistics.	944
891	relation extraction with directionality. <i>Bioinformat-</i>		
892	<i>ics</i> , 41(Supplement_1):i68–i76.	Mihir Parmar, Swaroop Mishra, Mor Geva, and Chitta	945
		Baral. 2023. Don’t blame the annotator: Bias al-	946
893	Nicolas Le Guillaume and Wilfried Thuiller. 2022. Tax-	ready starts in the annotation instructions . In <i>Pro-</i>	947
894	oNERD: deep neural models for the recognition of	<i>ceedings of the 17th Conference of the European</i>	948
895	taxonomic entities in the ecological and evolution-	<i>Chapter of the Association for Computational Lin-</i>	949
896	ary literature . <i>Methods in Ecology and Evolution</i> ,	<i>guistics</i> , pages 1779–1789, Dubrovnik, Croatia. As-	950
897	13(3):625–641.	sociation for Computational Linguistics.	951
898	Kyubum Lee, Maria Livia Famiglietti, Aoife McMa-	Hao Peng, Tianyu Gao, Xu Han, Yankai Lin, Peng	952
899	hon, Chih-Hsuan Wei, Jacqueline Ann Langdon	Li, Zhiyuan Liu, Maosong Sun, and Jie Zhou. 2020.	953
900	MacArthur, Sylvain Poux, Lionel Breuza, Alan	Learning from context or names? an empirical study	954
901	Bridge, Fiona Cunningham, Ioannis Xenarios, and	on neural relation extraction . In <i>Proceedings of the</i>	955
902	Zhiyong Lu. 2018. Scaling up data curation using	<i>2020 Conference on Empirical Methods in Natural</i>	956
903	deep learning: An application to literature triage in	<i>Language Processing (EMNLP)</i> , pages 3661–3672,	957
904	genomic variation resources . <i>PLOS Computational</i>	Online. Association for Computational Linguistics.	958
905	<i>Biology</i> , 14(8):e1006390.		
		Jorrit H Poelen, James D Simons, and Chris J Mungall.	959
906	Jens Lehmann, Daniel Gerber, Mohamed Morsey, and	2014. Global biotic interactions: An open infrastruc-	960
907	Axel-Cyrille Ngonga Ngomo. 2012. DeFacto – deep	ture to share and analyze species-interaction datasets.	961
908	fact validation. In <i>The Semantic Web – ISWC 2012</i> ,	<i>Ecological Informatics</i> , 24:148–159.	962
909	volume 7649 of <i>Lecture Notes in Computer Science</i> ,		
910	pages 312–327. Springer.	Adam Poliak, Jason Naradowsky, Aparajita Haldar,	963
		Rachel Rudinger, and Benjamin Van Durme. 2018.	964
911	Omer Levy, Minjoon Seo, Eunsol Choi, and Luke	Hypothesis only baselines in natural language infer-	965
912	Zettlemoyer. 2017. Zero-shot relation extraction via	ence . In <i>Proceedings of the Seventh Joint Confer-</i>	966
913	reading comprehension . In <i>Proceedings of the 21st</i>	<i>ence on Lexical and Computational Semantics</i> , pages	967
914	<i>Conference on Computational Natural Language</i>	180–191, New Orleans, Louisiana. Association for	968
915	<i>Learning (CoNLL 2017)</i> , pages 333–342, Vancouver,	Computational Linguistics.	969
916	Canada. Association for Computational Linguistics.		
		Qwen Team. 2025. Qwen3 technical report. <i>arXiv</i>	970
917	Yucheng Lu and Kazimier Smith. 2025. Feeding LLM	<i>preprint arXiv:2505.09388</i> .	971
918	annotations to BERT classifiers at your own risk.		
919	<i>arXiv preprint arXiv:2504.15432</i> .	Jonathan A. Rees and Karen Cranston. 2017. Auto-	972
		mated assembly of a reference taxonomy for phylo-	973
920	Ling Luo, Po-Ting Lai, Chih-Hsuan Wei, Cecilia N	genetic data synthesis . <i>Biodiversity Data Journal</i> ,	974
921	Arighi, and Zhiyong Lu. 2022. BioRED: a rich	5:e12581.	975
922	biomedical relation extraction dataset. <i>Briefings in</i>		
923	<i>Bioinformatics</i> , 23(5):bbac282.	Sebastian Riedel, Limin Yao, and Andrew McCallum.	976
		2010. Modeling relations and their mentions with-	977
924	Quinn McNemar. 1947. Note on the sampling error	out labeled text . In <i>Machine Learning and Knowl-</i>	978
925	of the difference between correlated proportions or	<i>edge Discovery in Databases (ECML PKDD 2010)</i> ,	979
926	percentages. <i>Psychometrika</i> , 12(2):153–157.	Lecture Notes in Computer Science, pages 148–163,	980
		Barcelona, Spain. Springer Berlin Heidelberg.	981
927	Mike Mintz, Steven Bills, Rion Snow, and Dan Juraf-	Raul Rodriguez-Esteban, Ivan Iossifov, and Andrey	982
928	sky. 2009. Distant supervision for relation extrac-	Rzhetsky. 2006. Imitating manual curation of text-	983
929	tion without labeled data. In <i>Proceedings of ACL-</i>	mined facts in biomedicine . <i>PLoS Computational</i>	984
930	<i>IJCNLP</i> , pages 1003–1011.	<i>Biology</i> , 2(9):e118.	985

986	Shachar Rosenman, Alon Jacovi, and Yoav Goldberg.	Peter West, Chandra Bhagavatula, Jack Hessel, Jena D.	1043
987	2020. Exposing shallow heuristics of relation ex-	Hwang, Liwei Jiang, Ronan Le Bras, Ximing Lu,	1044
988	traction models with challenge data . In <i>Proceed-</i>	Sean Welleck, and Yejin Choi. 2022. Symbolic	1045
989	<i>ings of the 2020 Conference on Empirical Methods</i>	knowledge distillation: from general language mod-	1046
990	<i>in Natural Language Processing (EMNLP)</i> , pages	els to commonsense models . In <i>Proceedings of the</i>	1047
991	3702–3710, Online. Association for Computational	<i>2022 Conference of the North American Chapter of</i>	1048
992	Linguistics.	<i>the Association for Computational Linguistics: Hu-</i>	1049
993	Patrick Ruch, Emilie Pasche, Pierre-André Michel, and	<i>man Language Technologies</i> , pages 4602–4625, Seat-	1050
994	Déborah Caucheteur. 2024. BiotXplorer: Navigat-	tle, United States. Association for Computational	1051
995	ing evidence-based biotic interactions . <i>Biodiversity</i>	Linguistics.	1052
996	<i>Information Science and Standards</i> , 8:e135453.		
997	Oscar Sainz, Oier Lopez de Lacalle, Gorka Labaka,	Michael Wiegand, Josef Ruppenhofer, and Thomas	1053
998	Ander Barrena, and Eneko Agirre. 2021. Label ver-	Kleinbauer. 2019. Detection of abusive language:	1054
999	balization and entailment for effective zero and few-	the problem of biased datasets . In <i>Proceedings of the</i>	1055
1000	shot relation extraction . In <i>Proceedings of the 2021</i>	<i>2019 Conference of the North American Chapter of</i>	1056
1001	<i>Conference on Empirical Methods in Natural Lan-</i>	<i>the Association for Computational Linguistics: Hu-</i>	1057
1002	<i>guage Processing</i> , pages 1199–1212, Online and	<i>man Language Technologies, Volume 1 (Long and</i>	1058
1003	Punta Cana, Dominican Republic. Association for	<i>Short Papers)</i> , pages 602–608, Minneapolis, Min-	1059
1004	Computational Linguistics.	nesota. Association for Computational Linguistics.	1060
1005	Tal Schuster, Darsh Shah, Yun Jie Serene Yeo, Daniel	Shanchan Wu and Yifan He. 2019. Enriching pre-	1061
1006	Roberto Filizzola Ortiz, Enrico Santus, and Regina	trained language model with entity information for	1062
1007	Barzilay. 2019. Towards debiasing fact verification	relation classification . In <i>Proceedings of the 28th</i>	1063
1008	models . In <i>Proceedings of the 2019 Conference on</i>	<i>ACM International Conference on Information and</i>	1064
1009	<i>Empirical Methods in Natural Language Processing</i>	<i>and Knowledge Management (CIKM)</i> , pages 2361–2364,	1065
1010	<i>and the 9th International Joint Conference on Natu-</i>	Beijing, China. Association for Computing Machin-	1066
1011	<i>ral Language Processing (EMNLP-IJCNLP)</i> , pages	ery.	1067
1012	3419–3425, Hong Kong, China. Association for Com-		
1013	putational Linguistics.	Michihiro Yasunaga, Jure Leskovec, and Percy Liang.	1068
1014	Akshay Swaminathan, Ivan Lopez, William Wang,	2022. LinkBERT: Pretraining language models with	1069
1015	Ujwal Srivastava, Edward Tran, Aarohi Bhargava-	document links. In <i>Proceedings of the 60th Annual</i>	1070
1016	Shah, Janet Y. Wu, Alexander L. Ren, Kaitlin Caoili,	<i>Meeting of the Association for Computational Lin-</i>	1071
1017	Brandon Bui, Layth Alkhani, Susan Lee, Nathan	<i>guistics (Volume 1: Long Papers)</i> , pages 8003–8016,	1072
1018	Mohit, Noel Seo, Nicholas Macedo, Winson Cheng,	Dublin, Ireland. Association for Computational Lin-	1073
1019	Charles Liu, Reena Thomas, Jonathan H. Chen,	guistics.	1074
1020	and Olivier Gevaert. 2024. Selective prediction	Yuhao Zhang, Victor Zhong, Danqi Chen, Gabor Angeli,	1075
1021	for extracting unstructured clinical data . <i>Journal</i>	and Christopher D. Manning. 2017. Position-aware	1076
1022	<i>of the American Medical Informatics Association</i> ,	attention and supervised data improve slot filling .	1077
1023	31(1):188–197.	In <i>Proceedings of the 2017 Conference on Empiri-</i>	1078
1024	Bruno Taillé, Vincent Guigue, Geoffrey Scuttheeten,	<i>cal Methods in Natural Language Processing</i> , pages	1079
1025	and Patrick Gallinari. 2020. Let’s stop incorrect	35–45, Copenhagen, Denmark. Association for Com-	1080
1026	comparisons in end-to-end relation extraction! In	putational Linguistics.	1081
1027	<i>Proceedings of the 2020 Conference on Empirical</i>	Wenxuan Zhou and Muhao Chen. 2022. An improved	1082
1028	<i>Methods in Natural Language Processing (EMNLP)</i> ,	baseline for sentence-level relation extraction . In <i>Pro-</i>	1083
1029	pages 3689–3701. Association for Computational	<i>ceedings of the 2nd Conference of the Asia-Pacific</i>	1084
1030	Linguistics.	<i>Chapter of the Association for Computational Lin-</i>	1085
1031	James Thorne, Andreas Vlachos, Christos	<i>guistics and the 12th International Joint Conference</i>	1086
1032	Christodoulopoulos, and Arpit Mittal. 2018.	<i>on Natural Language Processing (Volume 2: Short</i>	1087
1033	FEVER: a large-scale dataset for fact extraction and	<i>Papers)</i> , pages 161–168, Online only. Association for	1088
1034	VERification. In <i>Proceedings of NAACL-HLT</i> , pages	Computational Linguistics.	1089
1035	809–819.	Heng-Xing Zou, Xiaohao Yang, Thabassum H. Haja-	1090
1036	David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu	maideen, Olivia J. Stein, Roxanne S. Beltran, Ben-	1091
1037	Wang, Madeleine van Zuylen, Arman Cohan, and	jamin G. Freeman, Mark Lindquist, Eliot T. Miller,	1092
1038	Hannaneh Hajishirzi. 2020. Fact or fiction: Verifying	Summer Mengarelli, Charlotte M. Probst, Fernanda S.	1093
1039	scientific claims . In <i>Proceedings of the 2020 Con-</i>	Valdovinos, Derek Van Berkel, Phoebe L. Zarnetske,	1094
1040	<i>ference on Empirical Methods in Natural Language</i>	Brian C. Weeks, and Kai Zhu. 2026. Large language	1095
1041	<i>Processing (EMNLP)</i> , pages 7534–7550, Online. As-	models unlock the ecology of species interactions .	1096
1042	sociation for Computational Linguistics.	<i>bioRxiv</i> . Preprint, posted 10 February 2026.	1097

seed	AUPRC			F1, $\tau = 0.5$		
	sent.	pair	triple	sent.	pair	triple
1	.858	.922	.912	.632	.790	.842
2	.851	.921	.905	.698	.758	.847
3	.845	.912	.909	.663	.788	.834
mean	.851	.918	.909	.664	.779	.841
sd	.006	.005	.004	.033	.018	.006

Table 7: Per-seed figures on the 437 biodiversity rows, order-free scoring. The body reports three-checkpoint ensembles at block-held-out thresholds; a fixed $\tau = 0.5$ is shown here because the three models are calibrated differently, and a single shared threshold would favour one of them arbitrarily.

Model	Target	AUPRC	F1	AUPRC, taxa	
				≤ 2	≥ 3
Sentence	verdict	0.851 $\pm$.006	0.775	0.876	0.846
	$P(\text{YES})$	0.857 $\pm$.009	0.769	0.885	0.850
Pair	verdict	0.918 $\pm$.005	0.871	0.918	0.930
	$P(\text{YES})$	0.920 $\pm$.006	0.873	0.924	0.927
Triple	verdict	0.909 $\pm$.004	0.852	0.879	0.930
	$P(\text{YES})$	0.919 $\pm$.004	0.873	0.908	0.929

Table 8: The three arms trained on the teacher’s verdicts (the body’s models) and on its probabilities, recipe and seeds unchanged; columns as in Tables 1 and 2. Per-seed AUPRC with soft targets: .847/.864/.861, .927/.917/.917 and .920/.921/.915. Against the verdict-trained model of the same format, the soft-target model fixes 14 items and breaks 13 (sentence), 13 and 14 (pair), 19 and 7 (triple; McNemar $p = 1.0, 1.0, 0.03$).

A Per-seed results

B Soft-label distillation

The students learn from the teacher’s verdicts. We re-queried the teacher for $P(\text{YES})$ on all 48,338 training rows (same prompt and inputs; its verdicts agree with the corpus labels on 99.7% of rows, and 24% of the probabilities lie between 0.05 and 0.95) and trained the three arms again against these soft targets with a soft cross-entropy, selecting checkpoints on the same hard-labelled development split. It is the only distillation variant we tried, scored once (Table 8). The sentence-level and pair-conditioned models do not move. The triple-conditioned model gains 0.010 AUPRC, on every seed, and 0.021 F1 ($p = 0.03$ before correcting for three comparisons), which brings it level with the pair model (0.919 against 0.920; 24 items fixed, 22 broken, $p = 0.88$) rather than past it. The comparison is unchanged: the pair query lifts AUPRC from 0.857 to 0.920 (80 fixed, 29 broken,

Encoder	BioRED, BC8 test			Biodiversity, 437			
	sent.	pair	EP F1	sent.	pair	triple	pair F1
BiomedBERT-base	0.738	0.843	57.5/67.1	0.851	0.918	0.909	0.871
BiomedBERT-large	0.741	0.843	57.6/67.4	—	—	—	—
BioLinkBERT-base	0.743	0.855	57.9/67.9	0.869	0.921	0.922	0.861
BioLinkBERT-large	0.737	0.848	56.1/ 68.2	0.858	0.927	0.924	0.885

Table 9: The arms under the encoders that were scored on test, recipe otherwise that of §3: AUPRC (per-seed means, three seeds), EP F1 (sentence-level/pair-conditioned) and block-held-out F1. BiomedBERT-large was not trained on the biodiversity arms (Appendix K). Against the base encoder’s model of the same format, BioLinkBERT’s pair model fixes 621 BioRED items and breaks 537 (McNemar $p = 0.015$) and on ours 20 and 26 ($p = 0.46$); its sentence-level model on ours, 43 and 19 ($p = 0.003$). No BioLinkBERT-large arm differs significantly from the base encoder’s on ours ($p \geq 0.15$).

$p = 1.7 \times 10^{-6}$), by +0.077 on passages naming three or more taxa and +0.039 on the rest. Nor do probabilities close the pair model’s distance to its teacher’s 0.946 (0.026 AUPRC, against 0.027 with verdicts). The body keeps the verdict-trained models, on which every other analysis was run.

C Encoders

Three of the seven alternative recipes of §4.3 beat the base recipe by 0.01 development AUPRC or more on both splits (seed 1, pair model) and were scored once on the test sets: BiomedBERT-large at 2×10^{-5} (Appendix K), and BioLinkBERT, pre-trained on PubMed with citation links (Yasunaga et al., 2022), in its base size under the base recipe and in its large size (below). BioLinkBERT-base’s development AUPRC is 0.910 for the BioRED pair model over three seeds, against 0.892 for the base encoder and 0.905 for BiomedBERT-large, and 0.852 for the biodiversity pair model on seed 1, then 0.843 and 0.833 on seeds 2 and 3. Its BioRED pair model gains on every seed (AUPRC 0.853–0.858 against 0.840–0.847; EP F1 67.7–68.0 against 66.8–67.5), the best sentence-scope AUPRC we have on BioRED (Table 9). On our benchmark its pair model does not gain (0.921 ± 0.005 against 0.918 ± 0.005 ; F1 0.861 against 0.871) while its sentence-level model does (0.869, F1 0.803 against 0.775), so the formulation gap narrows to +0.052 AUPRC (54 items fixed, 31 broken, $p = 0.017$), still concentrated on passages naming three or more taxa (+0.070 against +0.019). On BioRED it widens, to +0.113.

BioLinkBERT-large, at 2×10^{-5} for three epochs like BiomedBERT-large, passed the same screen (development AUPRC 0.907 on BioRED and 0.869 on ours, seed 1; three-seed means 0.903 and 0.862). Its pair models gain a little on both benchmarks: on BioRED, AUPRC 0.848 ± 0.009 and EP F1 68.2 ± 0.7 (seeds 67.7–69.0, each above every base seed); on ours, 0.927 ± 0.004 and F1 0.885. Neither gain is significant: against the base encoder’s pair model it fixes 652 BioRED items and breaks 653 ($p = 1.0$), and on ours 27 and 21 ($p = 0.47$). Its other arms, trained afterwards and scored once, complete the comparison: the sentence-level models reach 0.858 on ours and 0.737 on BioRED (EP F1 56.1), the triple-conditioned model 0.924, and the pair query again wins, by +0.069 on ours (74 items fixed, 28 broken, $p = 8 \times 10^{-6}$; +0.083 on passages naming three or more taxa, +0.041 on the rest) and by +0.111 on BioRED. We keep the base encoder: BioLinkBERT-base does not carry to the pair-conditioned model the pipeline runs, and BioLinkBERT-large’s gains are not significant on either benchmark.

D Recall side

REJECT50 asks whether the candidates the retrieval pipeline discarded contain recoverable positives, and it is too small to rank the models: nine positives, so one item moves recall by eleven points. At block-held-out thresholds the sentence-level model recovers 7 of the 9 by re-admitting 15 of the 32 correct rejections (F1 0.452), the pair-conditioned model recovers 4 while re-admitting 7 (0.400), and the triple-conditioned model recovers 6 while re-admitting 5 (0.600). We read only the qualitative point: the discarded candidates do contain recoverable interactions, and a model that reads the candidate finds them at a lower cost in wrong re-admissions than one that reads only the passage.

E Teacher prompts

The three questions put to the 32B teacher in Table 1, verbatim, at temperature 0 with thinking disabled; the score is the probability of YES among the YES/NO alternatives of the first answer token. The triple question is the prompt that labelled the training corpus. On BioRED the sentence and pair questions name biomedical entities and relation types in place of organisms and biotic interactions (scripts/llm_baseline.py). Lines

are wrapped for the page.

Sentence question.

You check sentences retrieved from the scientific literature for a biotic-interaction database.

Sentence: <passage>

Answer YES only if the sentence asserts a direct biological interaction between organisms (for example predation, parasitism, infection, herbivory, pollination, mutualism, symbiosis or disease transmission), and the interaction is not negated.

Reply with exactly one word: YES or NO.

Pair question.

You verify candidate biotic-interaction pairs extracted from scientific literature.

Sentence: <passage>

Candidate pair:

Entity 1: "<taxon 1>"

Entity 2: "<taxon 2>"

Answer YES only if ALL of the following hold:

- Both entities are genuinely the organisms referred to by those surface strings in this sentence (not a gene, protein, chemical, author name, or a different species).
- The sentence asserts a direct biological interaction between these two organisms (not merely a co-mention, and not an interaction each has with some third organism), and the interaction is not negated.

Reply with exactly one word: YES or NO.

Triple question (the labelling prompt).

You verify candidate biotic-interaction triples extracted from scientific literature.

Sentence: <passage>

Candidate triple:

Entity 1: "<taxon 1>"

Relation: "<relation>"

Entity 2: "<taxon 2>"

Answer YES only if ALL of the following hold:

- Both entities are genuinely the organisms referred to by those surface strings in this sentence (not a gene, protein, chemical, author name, or a different species).
- The sentence asserts a direct biological interaction between these two organisms (not merely a co-mention, and not an interaction each has with some third organism).
- The stated relation is the correct type of interaction between them, in the stated direction, and is not negated.

Reply with exactly one word: YES or NO.

F Zero-shot scaling

Table 10 repeats the zero-shot comparison across model sizes. At 0.6B the three questions sit close

Qwen3	Biodiversity, 437			BioRED, 3,000	
	sent.	pair	triple	sent.	pair
0.6B	0.640	0.644	0.638	0.522	0.523
1.7B	0.753	0.794	0.828	0.533	0.648
4B	0.740	0.864	0.903	0.546	0.745
8B	0.802	0.898	0.936	0.524	0.800
14B	0.797	0.920	0.949	0.555	0.805
30B-A3B	0.777	0.902	0.933	0.571	0.792
32B (teacher)	0.789	0.943	0.946	0.566	0.813
Qwen3.5-122B	0.773	0.968	0.965	0.490 [†]	0.811 [†]
trained, same rows	0.855	0.923	0.909	0.757	0.860

Table 10: AUPRC of zero-shot Qwen3 (Qwen Team, 2025) by size, asked the three questions of Appendix E (4-bit builds served locally, thinking disabled; the triple question in both argument orders, maxed). 30B-A3B is a mixture of experts with 3B active parameters; Qwen3.5-122B, a later release, is a mixture of experts with 8 of 256 experts active per token, run partly on the CPU. Base rates 0.563 and 0.476. Last row: the trained three-checkpoint ensembles on the same rows. [†]On the first 1,000 candidates of the sample (base rate 0.442), where the 32B model scores 0.527 and 0.796 and the trained models 0.714 and 0.841.

to the base rate (0.64 against 0.563; 0.52 against 0.476): at that size the model barely does the task, whatever it is asked. From 1.7B on, the pair question is ahead at every size on both benchmarks: by +0.040 AUPRC at 1.7B and by +0.10 to +0.20 from 4B on (+0.11, then +0.20 to +0.32, on BioRED). The sentence question gains little with size (0.74–0.80 on ours from 4B on; 0.49–0.57 on BioRED throughout, near its base rate), so scale mostly improves the answer to the right question. The lead is not monotonic in size among the Qwen3 models, and it is largest for the 122B model: +0.195 on ours and +0.321 on its BioRED subsample, where the 32B model’s is +0.270. Asked about the pair, the 122B model ranks our benchmark better than any trained model (0.968; the pair-conditioned student 0.918), which suggests that the headroom left on it lies in a stronger teacher rather than in a different student recipe (Appendices B and C); asked the sentence-level question, it scores 0.773.

G Escalating uncertain candidates to the teacher

The deployed model is cheap and recall-oriented, and the zero-shot teacher, asked about the pair, is precise but slow (Appendix F). Table 11 combines them: the model decides every candidate, and those

Band	LLM calls	Qwen3-32B			Qwen3.5-122B		
		P	R	F1	P	R	F1
0%	0	0.869	0.947	0.907	0.869	0.947	0.907
10%	44	0.888	0.939	0.913	0.892	0.943	0.917
20%	87	0.900	0.919	0.909	0.912	0.923	0.917
30%	131	0.929	0.911	0.920	0.941	0.911	0.926
50%	218	0.946	0.850	0.895	0.958	0.837	0.894
100%	437	0.941	0.772	0.848	0.969	0.760	0.852

Table 11: The deployed model with its candidate rules keeps its own decision except on the band of candidates whose score lies nearest its threshold, which take the zero-shot teacher’s answer to the pair question; 437 rows. A rule rejection is final. Band 0% is the deployed model alone, 100% the teacher alone (with the rules).

it is least sure of — a fixed fraction, nearest its threshold — are re-decided by the teacher. Escalating 30% to Qwen3.5-122B raises precision from 0.869 to 0.941 at recall 0.911, removing 21 of the 35 false positives; the model alone reaches that precision only at $\tau = 0.99$, where recall falls to 0.740. Against the model alone the decisions are not significantly better on these 437 rows (25 items fixed, 13 broken, McNemar $p = 0.07$; with Qwen3-32B, 22 and 13, $p = 0.18$): the gain is in the operating point, more recall at a given precision, not in accuracy at the model’s own threshold. The band is a dial, reported at six settings and chosen at none of them. Its cost is set by the teacher’s speed: escalating 30% of the 175,588-candidate pool means about 53,000 calls, under a day on one A100 for Qwen3.5-122B (1.4 s per call with part of it on the CPU) and less for the 32B.

H Error examples

I Candidate rules in detail

The eight rules of §6 are applied to the candidate before the verifier decides, and each targets an error mechanism of §5. *same_organism* (apposition of a common and a scientific name, or identical forms), *own_clade* (an adjacent clade annotation such as “(Acari: Prostigmata)”), and *authority* (a taxonomic author parsed as a taxon) detect the same-organism and entity classes. *non_biotic_relation* rejects phylogenetic, comparative and biogeographic relation terms; *negated* rejects the two negation forms whose scope provably contains one of the pair (“non-*s* hosts”, “no infections with ... *s*”); *not_an_organism* rejects organs and syndromes; *pathogen_modifier* rejects an adjective that only ever occurs inside a pathogen’s name; and

Mechanism	Candidate triple	Passage (abridged)
Pair-binding	(<i>T. gondii</i> , parasite of, swine)	“The entero-pathogenic <i>Cystoisospora suis</i> , a coccidian parasite of swine and close relative to <i>Toxoplasma gondii</i> . . .”
Co-listing	(katyids, prey, beetles)	“... remains were from insects ... (e.g. beetles , dragonflies, cockroaches, and female katyids), indicating that eavesdropping on prey signals is not the only prey-finding strategy used by this <i>bat</i> .”
Negation	(wheat, hosts, <i>F. acuminatum</i>)	“... <i>F. acuminatum</i> in New Zealand was found only on non-wheat hosts .”
Same organism	(Chinese knotweed, pollinators, <i>P. chinensis</i>)	“... potential pollinators of Chinese knotweed (<i>Persicaria chinensis</i>).”
Entity	(Siberia , migrating, birds)	“... a portion of birds migrating from Siberia have common winter nesting areas.”

Table 12: One example of each mechanism among the pair-conditioned model’s false positives. The relation’s true argument is italicised where it is not the queried taxon; in the entity case a place name was matched as a taxon.

System	P	R	F1	fixed	broken
Pair-conditioned	0.849	0.894	0.871	—	—
+ 7 rules	0.858	0.886	0.872	3	2
+ 8 rules	0.867	0.874	0.870	6	5
Triple-conditioned	0.820	0.886	0.852	—	—
+ 7 rules	0.840	0.878	0.859	7	2
+ 8 rules	0.862	0.866	0.864	14	5

Table 13: Deterministic candidate rules applied before the model’s decision, all 437 rows, block-held-out thresholds. “Fixed” and “broken” are relative to the same model without rules. McNemar: $p = 1.00$ and 1.00 on the pair model, 0.18 and 0.067 on the triple model.

co_listed rejects two taxa that are conjuncts of one dependency-parsed list not governed by an interaction noun (“interactions between *A* and *B*” is kept).

How they were validated, and what that licenses. No rule has a parameter fitted to evaluation labels, but rules can still be fitted by *choice*, and four of these — *non_biotic_relation*, *negated*, *not_an_organism* and *pathogen_modifier* — were written after reading the false positives of §5. To keep their selection from simply memorising the benchmark, every rule had to pass a criterion fixed before it was evaluated: on the 42,747 training rows, at least 85% of the candidates it rejects must be teacher-negative, matching the teacher’s own accuracy. Two first drafts failed it and were narrowed on training data alone — a general negation cue fired on “without affecting the host”, and a role-noun list rejected “viruses” and “fungi”, which are legitimate coarse taxa — and the final rules reject teacher-negatives at 0.88 to 1.00 (*co_listed*, which needs a dependency parse, on a fixed 6,000-row sample). The rules were then frozen, recorded by

hash, and evaluated once.

J Where the reformulation stops

Entity disambiguation, and a gate that cannot fix it. The natural remedy for entity errors is to gate candidates on taxonomic validity. Against the Open Tree Taxonomy (Rees and Cranston, 2017) such a gate catches 1 of the 35 unsupported BIOTX100 candidates and discards 5 of the 65 supported ones. The reason is diagnostic: *Illiger*, *Dikarya*, *copia* and *sulla* — the surnames, clades, transposon families and vernacular names behind the entity failures — all *resolve* in the taxonomy. The failure is disambiguation, of the kind species-name normalisers address (Gerner et al., 2010), not validity; and the gate’s casualties are viruses, which are legitimate partners in this domain. The *authority* and *same_organism* rules (§6) target the disambiguation directly and are the better instrument.

Direction needs its own gold. Expert curation of the benchmark marks reversed triples as correct — we verified two by hand, (*Breinlia*, host of, *Rattus*) and (*Demidospermus*, parasitized by, *A. inermis*) — so no directional number on the benchmark would measure the model rather than the convention, which is why direction is scored on its own annotation (§7). The supervision itself is uneven in a revealing way: among the 5,591 reversed training rows, the rate at which the teacher rejects a reversal is governed by the relation’s *surface string*. Direction-encoding participles and prepositional forms are rejected at 0.51–0.70 (*affecting* 0.70, *affected by* 0.61, *parasite of* 0.57), bare relational nouns at 0.04–0.10 (*host of* 0.08, *pollinators* 0.04). Whether supervision encodes order at all depends on which string retrieval happened to match.

K Negative results and cost

Intervention	Benefit	Cost
Taxon-validity gate	1 of 35 unsupported caught	−5/65 supported; rejects viruses
Self-distillation	none	−0.001 P, −0.008 F1
6-layer compression	1.9× CPU throughput	−0.088 precision
Corpus signal-matching	none	−0.059 AUPRC
Rule-family training data	none	AUPRC +0.005 to −0.037
BiomedBERT-large	+0.013 AUPRC (BioRED pair)	dev none on test; 3× parameters

Table 14: Six plausible design choices that measurably do not pay.

AUPRC 0.787 and 0.791 against 0.786, inside the spread of the three-epoch seeds (0.765–0.786), so neither was trained further or scored on the test set.

Table 14 collects interventions that did not pay, because each is a plausible first thing to try. Self-distillation at equal capacity is null, suggesting the teacher’s knowledge is already saturated in the student. Compression to six layers buys throughput the pipeline does not need for 0.088 precision. Reweighting the corpus toward the deployment signal distribution cost 0.059 AUPRC. Most instructive is the last row: before the candidate rules existed, we turned the same mechanisms into *training data* — synthetic negatives for the same-organism, own-clade and authority classes, and in a second variant for co-listing — and retrained. The first variant left AUPRC unchanged within seed noise and the second cost 0.037. Mechanisms that are properties of the candidate are better answered by a rule than taught to a classifier. Deltas in the first four rows were measured under an earlier version of the evaluation and are given as directions. The last row is one of the three recipes of §4.3 that passed the development screen (the others are in Appendix C): BiomedBERT-large raises the BioRED pair model’s development AUPRC from 0.892 to 0.905 (three-seed means), but on the BC8 test set the pair model scores 0.843 ± 0.005 against 0.843 ± 0.004 , F1 0.814 against 0.817, EP F1 67.4 ± 0.4 against 67.1 ± 0.3 , and the sentence-level model 0.741 against 0.738; the pair query’s gain is unchanged (+0.103 against +0.105 AUPRC). It was not trained on all biodiversity arms, since its BioRED test result could not change the choice. For the whole-abstract BioRED model of §4.2 we also tried, on development data only (seed 1), eight epochs instead of three, alone and with each concept’s BioRED type in the query: development